

ROB 501 - Mathematics for Robotics

HW #7

Due 23:59 on Thursday, Oct. 17, 2024
To be submitted on Canvas

- **Remark A)** One of the main focuses of this HW set is Recursive Least Squares (RLS), in other words, least squares estimation that can be implemented in a real-time environment. This is a good way to prepare for the Kalman Filter. The HW set looks long because lots of details are given for each step in a problem, which should make the work go quickly.
- **Remark B)** To work the two problems on RLS, read the handouts `WeightedLeastSquares_RLS.pdf`, which is typeset and the handout `RecursiveLeastSquares_v02.pdf`, which is handwritten. These are under "Files/Handouts (Mostly Typeset by Grizzle on Individual Topics)" on Canvas. They are basically the same, but the handwritten one has a bit more detail. Read the handout and then implement the algorithm. This may be a bit tough on you, but when we go through it in lecture, you will be remarkably well prepared. The hints give you a LOT OF INFORMATION. Read them!

1. Prove, for a matrix $A \in \mathbb{R}^{m \times n}$, $\text{rank}(A) + \text{nullity}(A) = n$.

Remark: This is the *rank-nullity theorem* that we saw in lecture.

2. The symmetric matrix below has distinct e-values. Factor it as a product $O\Lambda O^\top$ where O is an orthogonal matrix and Λ is diagonal. It is OK to use MATLAB.

$$A = \begin{bmatrix} 2 & -1 & 0 \\ -1 & 2 & 1 \\ 0 & 1 & 2 \end{bmatrix}$$

3. The symmetric matrix below has repeated e-values $2, 2, -1$. In lecture, we stated that even with repeated e-values, we can still diagonalize a symmetric matrix using orthogonal matrices. The objective of the problem is to see why this is true by working a numerical example. We will follow the proof attached at the end of the HW set and factor A as a product $O\Lambda O^\top$, where O is an orthogonal matrix.

$$A = \begin{bmatrix} 1 & 0 & \sqrt{2} \\ 0 & 2 & 0 \\ \sqrt{2} & 0 & 0 \end{bmatrix}.$$

Each of the steps below is motivated by a step in the proof. **Suggestion:** Open a script file in MATLAB and execute each step of the problem. It will save time.

- (a) Verify that $v^1 = [0, 1, 0]^\top$ satisfies $Av^1 = 2v^1$, and thus v^1 is an e-vector corresponding to $\lambda = 2$.
- (b) Choose v^2 and v^3 such that $\{v^1, v^2, v^3\}$ is orthonormal, and verify that $V = [v^1 | v^2 | v^3]$ is an orthogonal matrix. In general, you would accomplish this by completing $\{v^1\}$ to a basis of \mathbb{R}^n and applying Gram Schmidt. Here, you can do it by inspection.

- (c) Form the matrix $V^\top AV$ and verify that it has the form

$$\begin{bmatrix} 2 & 0_{1 \times 2} \\ 0_{2 \times 1} & A_2 \end{bmatrix}$$

with A_2 symmetric.

- (d) Use MATLAB to compute the e-values and e-vectors of A_2 , and verify that they are 2, -1, and thus distinct¹. Find U_2 orthogonal such that

$$U_2^\top A_2 U_2$$

is diagonal. It is OK to use MATLAB for this.

- (e) Define the 3×3 matrix U by

$$U = \begin{bmatrix} 1 & 0_{1 \times 2} \\ 0_{2 \times 1} & U_2 \end{bmatrix}.$$

Verify that U is orthogonal.

- (f) Define $O = VU$ and verify that O is orthogonal.

- (g) **This is the only part you turn in:** Report what you get when you compute $O^\top AO$. If it is not diagonal, you have done something wrong.

4. Download the file `DataHW07_Prob4.mat` from the CANVAS MATLAB folder and load the file into your MATLAB workspace (See also the ReadMe.txt file). The data file provides “perturbed or noisy” data for the model $y_i = C_i x + e_i$, $1 \leq i \leq N$, where $N = 500$, $x \in \mathbb{R}^{100}$ and $y_i \in \mathbb{R}^3$. The data set contains the measured values y_i , the model matrices C_i , and the true value of x . The true value is given so that you can compare your estimated values to the true value. Of course, in real life, we would not have x available to us.

For $1 \leq k \leq N$, define $S_k = I_{3 \times 3}$ and as in lecture notes on Recursive Least Squares,

$$Y_k = \begin{bmatrix} y_1 \\ \vdots \\ y_k \end{bmatrix}, \quad A_k = \begin{bmatrix} C_1 \\ \vdots \\ C_k \end{bmatrix}, \quad R_k = \text{diag}[S_1, \dots, S_k] = I.$$

- (a) Find n such that A_k has at least $\dim(x) = 100$ independent columns for $k \geq n$. For each $n \leq k \leq N$, define

$$\hat{x}_k := \arg \min \|Y_k - A_k x\| = \arg \min \sqrt{(Y_k - A_k x)^\top R_k (Y_k - A_k x)}$$

but do not compute anything except n at this step.

- (b) For each $n \leq k \leq N$, compute \hat{x}_k in a batch process, that is,

$$\hat{x}_k = (A_k^\top R_k A_k)^{-1} A_k^\top R_k Y_k,$$

and, using the standard Euclidean norm, compute

$$E_k := \|\hat{x}_k - x\|.$$

¹If the e-values were repeated, you would find one e-vector, use it to build an orthonormal basis, and decompose A_2 to find A_3 symmetric that has dimension one less than A_2 , etc.

Make a plot of E_k versus k and turn it in. Put a clear title on your plot, such as “Norm error in \hat{x} using Batch Process”. Implementing the “for each $n \leq k \leq N$ ” will require a `for loop` or `while loop`. Use the `tic` and `toc` commands to determine how long it takes to compute your *entire set of estimates* and report this value. Either write it on your error plot by hand or place it there with a MATLAB command.

- (c) For each $n \leq k \leq N$, compute \hat{x}_k using the RLS (Recursive Least Squares) Algorithm. First implement it without using the Matrix Inversion Lemma. Turn in a plot of E_k versus k , and record on your plot the amount of time it takes to do your computations.
 - (d) For each $n \leq k \leq N$, compute \hat{x}_k once again using the RLS (Recursive Least Squares) Algorithm, but this time, implement it using the Matrix Inversion Lemma. Turn in a plot of E_k versus k , and record on your plot the amount of time it takes to do your computations. Note that this time you are numerically inverting a 3×3 matrix and then computing the inverse of the 100×100 matrix Q_k with the Matrix Inversion Lemma. This is the main point of the Matrix Inversion Lemma.
5. Download the file `DataHW07_Prob5.mat` from the CANVAS MATLAB folder and load the file into your MATLAB workspace. It provides “perturbed or noisy” data for the model $y_i = C_i x_i + e_i$, $1 \leq i \leq N$, where this time the “state” or “parameter” x that we are estimating is slowly “drifting” (means that it is slowly varying with time), which is why it has an index x_i . We will see that basic least squares does not work very well when x can drift. We will learn a way to fix it.

In this problem, $N = 500$, $x \in \mathbb{R}^{20}$ and $y_i \in \mathbb{R}^3$. The data set contains the measured values y_i , the model matrices C_i , and the true value of x_i . The true value is given so that you can compare your estimated values to the true value. As you know very well, in real life, we would not have x_i available to us.

- (a) Find n such that A_k has at least $\dim(x) = 20$ independent columns for $k \geq n$. For each $n \leq k \leq N$, define

$$\hat{x}_k := \arg \min \|Y_k - A_k x\|,$$

but do not compute anything except n at this step. You should find $n = 7$.

- (b) As in Prob. 4, use constant weights, with $S_k = I_{3 \times 3}$. For each $n \leq k \leq N$, compute \hat{x}_k (any method you wish) and compute $E_k := \|\hat{x}_k - x_k\|$. It does not matter how fast your MATLAB code is for the computation of \hat{x}_k because in this problem we will not record the time. Make a plot of E_k versus k and turn it in. Put a clear title on your plot. Note that the error gets pretty bad.
- (c) **The forgetting factor:** Let $0 < \lambda < 1$ (some number strictly between zero and one). A typical value for the *forgetting factor* might be $\lambda = 0.98$. The idea is to discount old measurements when we do the least squares problem. This is done by selecting at time k the weight matrices for $1 \leq i \leq k$ to be

$$S_i = \lambda^{(k-i)} I_{3 \times 3}.$$

With this choice, the $3k \times 3k$ weighting matrix R_k is given by

$$R_k = \text{diag}(\lambda^{k-1} I_3, \lambda^{k-2} I_3, \dots, \lambda I_3, I_3).$$

We see that the errors in older measurements are “discounted” by higher powers of λ , and thus the estimation process “exponentially forgets” them and “focuses” on the more recent measurements. It is important to note that at each step k , we are redefining the weights R_k so that errors in the newest measurements are penalized the most. This can be done recursively in our `for loop`, by

$$R_{k+1} = \begin{bmatrix} \lambda R_k & 0_{3k \times 3} \\ 0_{3 \times 3k} & I_{3 \times 3} \end{bmatrix}.$$

For each $n \leq k \leq N$, compute \hat{x}_k using the Batch Method. Turn in a plot of E_k versus k , and label your plot appropriately. You can use $\lambda = 0.98$ or you can tune the forgetting factor to see what works best. How can you resist playing with it once you have your code working? :)

- (d) For each $n \leq k \leq N$, compute \hat{x}_k now using the RLS (Recursive Least Squares) Algorithm, with forgetting factor. The algorithm (without using the Matrix Inversion Lemma) becomes

- **Initialization Step:** Set

$$Q_n := \sum_{i=1}^n C_i^\top \lambda^{n-i} C_i$$

$$\Gamma_n := \sum_{i=1}^n C_i^\top \lambda^{n-i} y_i$$

$$\hat{x}_n := (Q_n)^{-1} \Gamma_n$$

- **Recursion:** For $n \leq k < N$

$$Q_{k+1} := \lambda Q_k + C_{k+1}^\top C_{k+1}$$

$$K_{k+1} := (Q_{k+1})^{-1} C_{k+1}^\top$$

$$\hat{x}_{k+1} := \hat{x}_k + K_{k+1} (y_{k+1} - C_{k+1} \hat{x}_k)$$

- If you want the version with the Matrix Inversion Lemma, see the hints!
 - Turn in a plot of $E_k := \|\hat{x}_k - x_k\|$ versus k , and label your plot appropriately. To be clear, there are no λ 's in the computation of E_k ; we are just using the standard Euclidean norm to see how well we are doing in tracking x as it slowly drifts.
6. Classify each matrix as positive definite, positive semi-definite, or neither. In addition, if the matrix is either positive definite or positive semi-definite, find a square root. You may use MATLAB to factor a symmetric matrix as $\Lambda = O^\top P O$ or as $P = O \Lambda O^\top$.

(a) $\textcircled{\smiley} = \begin{bmatrix} 1 & 3 \\ 3 & 9 \end{bmatrix}$.

(b) $\textcircled{\frown} = \begin{bmatrix} 6 & 10 & 11 \\ 10 & 19 & 19 \\ 11 & 19 & 21 \end{bmatrix}$.

(c) $\textcircled{\smile} = \begin{bmatrix} 2 & 6 & 10 \\ 6 & 10 & 14 \\ 10 & 14 & 18 \end{bmatrix}$.

7. Use the results on Schur Complements to solve the following problems by hand:

(a) Determine if $\textcircled{\smiley} = \begin{bmatrix} 1 & 3 \\ 3 & 8 \end{bmatrix}$ is positive definite or not.

(b) Determine if $\textcircled{\frown} = \begin{bmatrix} 1 & 0 & 6 \\ 0 & 4 & 7 \\ 6 & 7 & 10 \end{bmatrix}$ is positive definite or not.

(c) Find the range of a such that the following matrix is positive definite: $\textcircled{\smile} = \begin{bmatrix} 1 & 2 & 6 \\ 2 & 5 & 7 \\ 6 & 7 & a \end{bmatrix}$

8. Find x of minimum norm that satisfies the equation

$$\begin{bmatrix} 1 & 3 & 2 \\ 3 & 8 & 4 \end{bmatrix} x = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

(a) Use the standard inner product on \mathbb{R}^3

(b) Use the inner product $\langle x, y \rangle = x^\top \begin{bmatrix} 5 & 1 & 9 \\ 1 & 2 & 1 \\ 9 & 1 & 17 \end{bmatrix} y$

Hints

Hints: Prob. 1 Recall the following facts:

- (a) The null space and range space are subspaces.
- (b) The nullity is $\dim \mathcal{N}(A)$ and the rank is $\dim \mathcal{R}(A)$ (dimension of null and range spaces of A).
- (c) Given a set of $p < n$ linearly independent vectors in \mathbb{R}^n , we can complete them to a basis for \mathbb{R}^n with $n - p$ more linearly independent vectors.

Recall also the definition of dimension as it relates to the basis, and the definition of a basis.

Hints: Prob. 3 The important point here is that when a matrix is symmetric, repeated e-values do not pose a problem as they do for a general square matrix. The last page of the HW gives a proof by induction.

- (a) **Base Step:** The first thing to note is that a 1×1 matrix can always be factored.
- (b) **Inductive Step:** The induction hypothesis is to assume that $(n - 1) \times (n - 1)$ symmetric matrices can be factored as $O\Lambda O^\top$ where O is an orthogonal matrix and Λ is diagonal.
- (c) **To show:** Next, you must show that the same is true for $n \times n$ symmetric matrices. The key step in the proof is to show that if A is symmetric and λ is an e-value, then there exists an orthogonal matrix P such that

$$P^\top A P = \begin{bmatrix} \lambda & 0_{1 \times (n-1)} \\ 0_{(n-1) \times 1} & B \end{bmatrix},$$

where B is symmetric and $(n - 1) \times (n - 1)$. The orthogonal matrix P is produced by using an e-vector associated with λ and the Gram-Schmidt process. Hence, if you care to understand the proof, it is within your means to do so.

Hints: Prob. 4 Recursive Least Squares (RLS)

(a) **Basic Version:**

- Initialization Step: Choose n such that Q_n is invertible (full rank)

$$\begin{aligned} Q_n &:= \sum_{i=1}^n C_i^\top S_i C_i \\ \Gamma_n &:= \sum_{i=1}^n C_i^\top S_i y_i \\ \hat{x}_n &:= (Q_n)^{-1} \Gamma_n \end{aligned}$$

- Recursion: For $n \leq k < N$

$$\begin{aligned} Q_{k+1} &:= Q_k + C_{k+1}^\top S_{k+1} C_{k+1} \\ K_{k+1} &:= (Q_{k+1})^{-1} C_{k+1}^\top S_{k+1} \\ \hat{x}_{k+1} &:= \hat{x}_k + K_{k+1} (y_{k+1} - C_{k+1} \hat{x}_k) \end{aligned}$$

(b) **Improved Version Using the Matrix Inversion Lemma:**

- Initialization Step: Choose n such that Q_n is invertible (full rank)

$$Q_n := \sum_{i=1}^n C_i^\top S_i C_i$$

$$P_n := (Q_n)^{-1}$$

$$\Gamma_n := \sum_{i=1}^n C_i^\top S_i y_i$$

$$\hat{x}_n := P_n \Gamma_n$$

- Recursion: For $n \leq k < N$

$$P_{k+1} = P_k - P_k C_{k+1}^\top [S_{k+1}^{-1} + C_{k+1} P_k C_{k+1}^\top]^{-1} C_{k+1} P_k.$$

$$K_{k+1} := P_{k+1} C_{k+1}^\top S_{k+1}$$

$$\hat{x}_{k+1} := \hat{x}_k + K_{k+1} (y_{k+1} - C_{k+1} \hat{x}_k)$$

- How to Derive the Riccati Equation? It comes from the Matrix Inversion Lemma

$$Q_{k+1} = Q_k + C_{k+1}^\top S_{k+1} C_{k+1}$$

$$Q_{k+1}^{-1} = (Q_k + C_{k+1}^\top S_{k+1} C_{k+1})^{-1}$$

$$= Q_k^{-1} - Q_k^{-1} C_{k+1}^\top [S_{k+1}^{-1} + C_{k+1} Q_k^{-1} C_{k+1}^\top]^{-1} C_{k+1} Q_k^{-1}$$

$$P_k := Q_k^{-1}$$

$$P_{k+1} = P_k - P_k C_{k+1}^\top [S_{k+1}^{-1} + C_{k+1} P_k C_{k+1}^\top]^{-1} C_{k+1} P_k.$$

- Jacopo Francesco Riccati (1676-1754) http://en.wikipedia.org/wiki/Jacopo_Riccati

Hints: Prob. 5

- (a) Rewrite $Q_{k+1} = \lambda Q_k + C_{k+1}^\top C_{k+1}$ as

$$\frac{1}{\lambda} Q_{k+1} = Q_k + C_{k+1}^\top \frac{1}{\lambda} C_{k+1}$$

Therefore

$$\lambda Q_{k+1}^{-1} = [Q_k + C_{k+1}^\top \frac{1}{\lambda} C_{k+1}]^{-1} \quad (*)$$

Using the Matrix Inversion Lemma, we have

$$\lambda Q_{k+1}^{-1} = Q_k^{-1} - Q_k^{-1} C_{k+1}^\top [\lambda I + C_{k+1} Q_k^{-1} C_{k+1}^\top]^{-1} C_{k+1} Q_k^{-1}.$$

and thus

$$Q_{k+1}^{-1} = \frac{1}{\lambda} Q_k^{-1} - \frac{1}{\lambda} Q_k^{-1} C_{k+1}^\top [\lambda I + C_{k+1} Q_k^{-1} C_{k+1}^\top]^{-1} C_{k+1} Q_k^{-1}.$$

If we define $P_k := Q_k^{-1}$, we obtain

$$P_{k+1} = \frac{1}{\lambda} P_k - \frac{1}{\lambda} P_k C_{k+1}^\top [\lambda I + C_{k+1} P_k C_{k+1}^\top]^{-1} C_{k+1} P_k.$$

- (b) If you are using your m-file for the Matrix Inversion Lemma, you can stop at (*), apply your function to get the inverse of $[Q_k + C_{k+1}^\top \frac{1}{\lambda} C_{k+1}]$, and then divide by the forgetting factor.
- (c) The recursion on \hat{x}_k is unchanged from the RLS algorithm without the forgetting factor. In case you want to see the derivation, the key formulas are:

$$\begin{aligned}
Q_k &:= \sum_{i=1}^k C_i^\top \lambda^{k-i} C_i \\
Q_k \hat{x}_k &:= \sum_{i=1}^k C_i^\top \lambda^{k-i} y_i \\
Q_{k+1} &= \sum_{i=1}^{k+1} C_i^\top \lambda^{k+1-i} C_i \\
&= \lambda Q_k + C_{k+1}^\top C_{k+1} \\
Q_{k+1} \hat{x}_{k+1} &= \sum_{i=1}^{k+1} C_i^\top \lambda^{k+1-i} y_i \\
&= \lambda \sum_{i=1}^k C_i^\top \lambda^{k-i} y_i + C_{k+1}^\top y_{k+1} \\
&= \lambda Q_k \hat{x}_k + C_{k+1}^\top y_{k+1} \\
\lambda Q_k &= Q_{k+1} - C_{k+1}^\top C_{k+1}
\end{aligned}$$

and thus, putting all of this together

$$\begin{aligned}
\hat{x}_{k+1} &= Q_{k+1}^{-1} [(Q_{k+1} - C_{k+1}^\top C_{k+1}) \hat{x}_k + C_{k+1}^\top y_{k+1}] \\
&= \hat{x}_k + Q_{k+1}^{-1} C_{k+1}^\top (y_{k+1} - C_{k+1} \hat{x}_k)
\end{aligned}$$

Hence, the only change is to the update formula for Q_{k+1} .

Hints: Prob. 6 Use the `help eig` command in MATLAB.

Hints: Prob. 7 Recall that for a symmetric matrix $M = \begin{bmatrix} A & B \\ B^\top & C \end{bmatrix}$ the following are equivalent:

- (a) $M \succ 0$
- (b) $A \succ 0$ and $C - B^\top A^{-1} B \succ 0$
- (c) $C \succ 0$ and $A - B C^{-1} B^\top \succ 0$


Hints: Prob. 8 This is an under determined system of equations and not an over determined system of equations.

Reply With Quote

Opalg
MHB Oldtimer

MHB Moderator

MHB Math Helper



Status: Offline

Join Date: Feb 2012

Location: Leeds, UK

Posts: 1,480

Thanks: 481 times

Thanked: 4510 times

Awards: 

DECEMBER 23RD, 2012, 12:21

#3

Originally Posted by **matqkks**

I have been trying to prove the following result:
If A is real symmetric matrix with an eigenvalue λ of multiplicity m then λ has m linearly independent e.vectors.
Is there a simple proof of this result?

This is a slight variation of Deveno's argument. I will assume you already know that the eigenvalues of a real symmetric matrix are all real.

Let A be an $n \times n$ real symmetric matrix, and assume as an inductive hypothesis that all $(n-1) \times (n-1)$ real symmetric matrices are diagonalisable. Let λ be an eigenvalue of A , with a normalised eigenvector x_1 . Using the Gram-Schmidt process, form an orthonormal basis $\{x_1, x_2, \dots, x_n\}$ with that eigenvector as its first element.

Let P be the $n \times n$ matrix whose columns are x_1, x_2, \dots, x_n , and denote by $T: \mathbb{R}^n \rightarrow \mathbb{R}^n$ the linear transformation whose matrix with respect to the standard basis is A . Then P is an orthogonal matrix ($P^T = P^{-1}$), and the matrix of T with respect to the basis $\{x_1, x_2, \dots, x_n\}$ is $P^T A P$. The (i, j) -element of that matrix is $(P^T A P)_{ij} = \langle A x_j, x_i \rangle$. In particular, the elements in the first column are

$$(P^T A P)_{i1} = \langle A x_1, x_i \rangle = \langle \lambda x_1, x_i \rangle = \begin{cases} \lambda & (i = 1) \\ 0 & (i > 1) \end{cases}$$

(because the vectors x_i are orthonormal). Thus the first column of $P^T A P$ has λ as its top element, and 0 for each of the other elements. Since $P^T A P$ is symmetric, the top row also consists of a λ followed by all zeros. Hence the matrix $P^T A P$ looks like this:

$$\begin{bmatrix} \lambda & 0 & \dots & 0 \\ 0 & & & \\ \vdots & & B & \\ 0 & & & \end{bmatrix},$$

where B is an $(n-1) \times (n-1)$ real symmetric matrix. By the inductive hypothesis, B is diagonalisable, so there is an orthogonal $(n-1) \times (n-1)$ matrix Q such that $Q^T B Q$ is

diagonal. Let

$$R = \begin{bmatrix} 1 & 0 & \dots & 0 \\ 0 & & & \\ \vdots & & Q & \\ 0 & & & \end{bmatrix}.$$

Then $R^T P^T A P R$ is diagonal, as required.

Was sich überhaupt sagen lässt, lässt sich klar sagen; und wovon man nicht reden kann, darüber muss man schweigen.

(Anything that can be said at all, can be said clearly; and whereof one cannot speak, thereon one must be silent.)

– Ludwig Wittgenstein's good advice for forum contributors, in *Tractatus Logico-Philosophicus*.

[Reply With Quote](#)

[▲ Go to First Post](#)

« Singular Value decomposition | An inverse of the adjoint »

Similar Threads

[SOLVED] Incorporate axis symmetric case

By dwsmith in forum Mathematics Software and Calculator Discussion

Replies: 0

Last Post: December 6th, 2012, 14:56

Inverse of a Symmetric Matrix

By OhMyMarkov in forum Linear and Abstract Algebra

Replies: 2

Last Post: October 6th, 2012, 11:28

[SOLVED] Eigenvalues

By Sudharaka in forum Pre-Algebra and Algebra

Replies: 2

Last Post: May 27th, 2012, 09:02

Matrix Theory...showing that matrix is Unitary

By cylv89 in forum Linear and Abstract Algebra

Replies: 3

Last Post: March 27th, 2012, 17:00

A and B are two symmetric matrices

By Yankel in forum Linear and Abstract Algebra

Replies: 4

Last Post: January 27th, 2012, 09:17

-- Math Help Boards

Have a small screen? Select Math Help Boards for Small Screens!

Powered by vBulletin Copyright © 2000 - 2012, Jelsoft Enterprises Ltd.

Search Engine Optimisation provided by DragonByte SEO v1.0.15 (Pro) - vBulletin Mods & Addons Copyright © 2014 DragonByte Technologies Ltd.

Feedback Buttons provided by Advanced Post Thanks / Like (Pro) - vBulletin Mods & Addons Copyright © 2014 DragonByte Technologies Ltd.

© 2012-2014 Math Help Boards

[FORUMS](#)

[RULES](#)

[POTW](#)

[CONTACT](#)

[TOP](#)

